

# Tutorial: Containerize a .NET Core app

01/09/2020 • 11 minutes to read •  +5

## In this article

[Prerequisites](#)

[Create .NET Core app](#)

[Publish .NET Core app](#)

[Create the Dockerfile](#)

[Create a container](#)

[Essential commands](#)

[Clean up resources](#)

[Next steps](#)

This tutorial teaches you how to build a Docker image that contains your .NET Core application. The image can be used to create containers for your local development environment, private cloud, or public cloud.

You'll learn to:

- ✓ Create and publish a simple .NET Core app
- ✓ Create and configure a Dockerfile for .NET Core
- ✓ Build a Docker image
- ✓ Create and run a Docker container

You'll understand the Docker container build and deploy tasks for a .NET Core application. The *Docker platform* uses the *Docker engine* to quickly build and package apps as *Docker images*. These images are written in the *Dockerfile* format to be deployed and run in a layered container.

### Tip

If you're working with an existing ASP.NET Core application, see the [Learn how to containerize an ASP.NET Core application](#) tutorial.

## Prerequisites

Install the following prerequisites:

- [.NET Core 3.1 SDK](#)


If you have .NET Core installed, use the `dotnet --info` command to determine which SDK you're using.

- [Docker Community Edition](#)


- A temporary working folder for the *Dockerfile* and .NET Core example app. In this tutorial, the name *docker-working* is used as the working folder.

## Create .NET Core app


You need a .NET Core app that the Docker container will run. Open your terminal, create a working folder if you haven't already, and enter it. In the working folder, run the following command to create a new project in a subdirectory named *app*:

.NET Core CLI	 Copy
<pre>dotnet new console -o app -n myapp</pre>	

Your folder tree will look like the following:

	 Copy
<pre>docker-working ├── app │   ├── myapp.csproj │   ├── Program.cs │   └── obj │       ├── myapp.csproj.nuget.cache │       ├── myapp.csproj.nuget.dgspec.json │       ├── myapp.csproj.nuget.g.props │       ├── myapp.csproj.nuget.g.targets │       └── project.assets.json</pre>	


The `dotnet new` command creates a new folder named *app* and generates a "Hello World" app. Enter the *app* folder and run the command `dotnet run`. You'll see the following output:

console	 Copy
<pre>&gt; dotnet run Hello World!</pre>	

The default template creates an app that prints to the terminal and then exits. For this tutorial, you'll use an app that loops indefinitely. Open the *Program.cs* file in a text editor. It should currently look like the following code:

C#	 Copy
<pre>using System;  namespace myapp {     class Program     {         static void Main(string[] args)         {             Console.WriteLine("Hello World!");         }     } }</pre>	

Replace the file with the following code that counts numbers every second:

C#	 Copy
<pre>using System;  namespace myapp {     class Program     {         static void Main(string[] args)         {             var counter = 0;             var max = args.Length != 0 ? Convert.ToInt32(args[0]) : -1;             while (max == -1    counter &lt; max)             {                 counter++;                 Console.WriteLine(\$"Counter: {counter}");                 System.Threading.Tasks.Task.Delay(1000).Wait();             }         }     } }</pre>	

Save the file and test the program again with `dotnet run`. Remember that this app runs indefinitely. Use the cancel command `CTRL+C` to stop it. You'll see the following output:

console	 Copy
---------	--

```
> dotnet run
Counter: 1
Counter: 2
Counter: 3
Counter: 4
^C
```

If you pass a number on the command line to the app, it will only count up to that amount and then exit. Try it with `dotnet run -- 5` to count to five.

### ⓘ Note

Any parameters after `--` are not passed to the `dotnet run` command and instead are passed to your application.

## Publish .NET Core app

Before you add your .NET Core app to the Docker image, publish it. You want to make sure that the container runs the published version of the app when it's started.

From the working folder, enter the *app* folder with the example source code and run the following command:

.NET Core CLI



```
dotnet publish -c Release
```

This command compiles your app to the *publish* folder. The path to the *publish* folder from the working folder should be `.\app\bin\Release\netcoreapp3.1\publish\`

From the *app* folder, get a directory listing of the publish folder to verify that the *myapp.dll* file was created.

console



```
> dir bin\Release\netcoreapp3.1\publish

        Directory: C:\docker-working\app\bin\Release\netcoreapp3.1\publish

01/09/2020  11:41 AM    <DIR>          .
01/09/2020  11:41 AM    <DIR>          ..
01/09/2020  11:41 AM                407 myapp.deps.json
01/09/2020  12:15 PM               4,608 myapp.dll
01/09/2020  12:15 PM            169,984 myapp.exe
```

01/09/2020	12:15 PM	736	myapp.pdb
01/09/2020	11:41 AM	154	myapp.runtimeconfig.json

If you're using Linux or macOS, use the `ls` command to get a directory listing and verify that the *myapp.dll* file was created.

bash	 Copy
me@DESKTOP:/docker-working/app\$ ls bin/Release/netcoreapp3.1/publish myapp.deps.json myapp.dll myapp.pdb myapp.runtimeconfig.json	

## Create the Dockerfile

The *Dockerfile* file is used by the `docker build` command to create a container image. This file is a text file named *Dockerfile* that doesn't have an extension.

In your terminal, navigate up a directory to the working folder you created at the start. Create a file named *Dockerfile* in your working folder and open it in a text editor. Depending on the type of application you're going to containerize, you'll choose either the ASP.NET Core runtime or the .NET Core runtime. When in doubt, choose the ASP.NET Core runtime, which includes the .NET Core runtime. This tutorial will use the ASP.NET Core runtime image, but the application created in the previous sections is an .NET Core application.

- ASP.NET Core runtime

Dockerfile	 Copy
FROM mcr.microsoft.com/dotnet/core/aspnet:3.1	

- .NET Core runtime

Dockerfile	 Copy
FROM mcr.microsoft.com/dotnet/core/runtime:3.1	

The `FROM` command tells Docker to pull down the image tagged **3.1** from the specified repository. Make sure that you pull the runtime version that matches the runtime targeted by your SDK. For example, the app created in the previous section used the .NET Core 3.1 SDK and the base image referred to in the *Dockerfile* is tagged with **3.1**.

Save the *Dockerfile* file. The directory structure of the working folder should look like the following. Some of the deeper-level files and folders have been cut to save space in

the article:

Copy

```
docker-working
├── Dockerfile
└── app
    ├── myapp.csproj
    ├── Program.cs
    ├── bin
    │   └── Release
    │       ├── netcoreapp3.1
    │       └── publish
    │           ├── myapp.deps.json
    │           ├── myapp.exe
    │           ├── myapp.dll
    │           ├── myapp.pdb
    │           └── myapp.runtimeconfig.json
    └── obj
```

From your terminal, run the following command:

console

Copy

```
docker build -t myimage -f Dockerfile .
```

Docker will process each line in the *Dockerfile*. The `.` in the `docker build` command tells Docker to use the current folder to find a *Dockerfile*. This command builds the image and creates a local repository named **myimage** that points to that image. After this command finishes, run `docker images` to see a list of images installed:

console

Copy

```
> docker images
REPOSITORY              TAG                IMAGE ID
CREATED                SIZE
myimage                 latest            38db0e-
b8f648                  4 weeks ago      346MB
mcr.microsoft.com/dotnet/core/aspnet  3.1              38db0e-
b8f648                  4 weeks ago      346MB
```

Notice that the two images share the same **IMAGE ID** value. The value is the same between both images because the only command in the *Dockerfile* was to base the new image on an existing image. Let's add two commands to the *Dockerfile*. Each command


creates a new image layer with the final command representing the image the **myimage** repository entry points to.

Dockerfile	 Copy
<pre>COPY app/bin/Release/netcoreapp3.1/publish/ app/  ENTRYPOINT ["dotnet", "app/myapp.dll"]</pre>	

The `COPY` command tells Docker to copy the specified folder on your computer to a folder in the container. In this example, the *publish* folder is copied to a folder named *app* in the container.

The next command, `ENTRYPOINT`, tells Docker to configure the container to run as an executable. When the container starts, the `ENTRYPOINT` command runs. When this command ends, the container will automatically stop.

From your terminal, run `docker build -t myimage -f Dockerfile .` and when that command finishes, run `docker images`.

console	 Copy
<pre>&gt; docker build -t myimage -f Dockerfile . Sending build context to Docker daemon 1.624MB Step 1/3 : FROM mcr.microsoft.com/dotnet/core/aspnet:3.1 ----&gt; 38db0eb8f648 Step 2/3 : COPY app/bin/Release/netcoreapp3.1/publish/ app/ ----&gt; 37873673e468 Step 3/3 : ENTRYPOINT ["dotnet", "app/myapp.dll"] ----&gt; Running in d8deb7b3aa9e Removing intermediate container d8deb7b3aa9e ----&gt; 0d602ca35c1d Successfully built 0d602ca35c1d Successfully tagged myimage:latest  &gt; docker images REPOSITORY              TAG                IMAGE ID CREATED                SIZE myimage                 latest            0d602- ca35c1d                4 seconds ago    346MB mcr.microsoft.com/dotnet/core/aspnet  3.1              38db0e- b8f648                4 weeks ago     346MB</pre>	

Each command in the *Dockerfile* generated a layer and created an **IMAGE ID**. The final **IMAGE ID** (yours will be different) is **ddcc6646461b** and next you'll create a container based on this image.

# Create a container

Now that you have an image that contains your app, you can create a container. You can create a container in two ways. First, create a new container that is stopped.

console	 Copy
<pre>&gt; docker create myimage ceda87b219a4e55e9ad5d833ee1a7ea4da21b5ea7ce5a7d08f3051152e784944</pre>	

The `docker create` command from above will create a container based on the **myimage** image. The output of that command shows you the **CONTAINER ID** (yours will be different) of the created container. To see a list of *all* containers, use the `docker ps -a` command:

console

Copy

> docker ps -a


CONTAINER ID	IMAGE	COMMAND	CRE-
ATED	STATUS	PORTS	NAMES
ceda87b219a4	myimage	"dotnet app/myapp.dll"	4
seconds ago	Created	gallant_lehmann	

## Manage the container

Each container is assigned a random name that you can use to refer to that container instance. For example, the container that was created automatically chose the name **gallant\_lehmann** (yours will be different) and that name can be used to start the container. You override the automatic name with a specific one by using the `docker create --name` parameter.

The following example uses the `docker start` command to start the container, and then uses the `docker ps` command to only show containers that are running:

console

 Copy

> docker start gallant\_lehmann  
gallant\_lehmann

> docker ps

CONTAINER ID	IMAGE	COMMAND	CRE-
ATED	STATUS	PORTS	NAMES
ceda87b219a4	myimage	"dotnet app/myapp.dll"	7
minutes ago	Up 8 seconds	gallant_lehmann	



Similarly, the `docker stop` command will stop the container. The following example uses the `docker stop` command to stop the container, and then uses the `docker ps` command to show that no containers are running:

console

Copy

> docker stop gallant\_lehmann  
gallant\_lehmann

> docker ps

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	

## Connect to a container


After a container is running, you can connect to it to see the output. Use the `docker start` and `docker attach` commands to start the container and peek at the output stream. In this example, the `CTRL + C` keystroke is used to detach from the running container. This keystroke may actually end the process in the container, which will stop the container. The `--sig-proxy=false` parameter ensures that `CTRL + C` won't stop the process in the container.

After you detach from the container, reattach to verify that it's still running and counting.


console					Copy
<pre>&gt; docker start gallant_lehmann gallant_lehmann  &gt; docker attach --sig-proxy=false gallant_lehmann Counter: 7 Counter: 8 Counter: 9 ^C  &gt; docker attach --sig-proxy=false gallant_lehmann Counter: 17 Counter: 18 Counter: 19 ^C</pre>					

## Delete a container

For the purposes of this article you don't want containers just hanging around doing nothing. Delete the container you previously created. If the container is running, stop it.

console	 Copy
<pre>&gt; docker stop gallant_lehmann</pre>	

The following example lists all containers. It then uses the `docker rm` command to delete the container, and then checks a second time for any running containers.

console	 Copy
<pre>&gt; docker ps -a CONTAINER ID        IMAGE               COMMAND             CRE- ATED              STATUS    PORTS      NAMES ceda87b219a4      myimage           "dotnet app/myapp.dll"  19 minutes ago      Exited                    gallant_lehmann  &gt; docker rm gallant_lehmann gallant_lehmann  &gt; docker ps -a CONTAINER ID        IMAGE               COMMAND             CREATED STATUS            PORTS              NAMES </pre>	

### Single run

Docker provides the `docker run` command to create and run the container as a single command. This command eliminates the need to run `docker create` and then `docker start`. You can also set this command to automatically delete the container when the container stops. For example, use `docker run -it --rm` to do two things, first, automatically use the current terminal to connect to the container, and then when the container finishes, remove it:

console	 Copy
<pre>&gt; docker run -it --rm myimage Counter: 1 Counter: 2 Counter: 3 Counter: 4 Counter: 5 ^C</pre>	

With `docker run -it`, the `CTRL + C` command will stop process that is running in the container, which in turn, stops the container. Since the `--rm` parameter was provided, the container is automatically deleted when the process is stopped. Verify that it doesn't exist:

console					Copy
<pre>&gt; docker ps -a</pre>					
CONTAINER ID	IMAGE	COMMAND	CRE-		
ATED	STATUS	PORTS	NAMES		

## Change the ENTRYPOINT

The `docker run` command also lets you modify the `ENTRYPOINT` command from the *Dockerfile* and run something else, but only for that container. For example, use the following command to run `bash` or `cmd.exe`. Edit the command as necessary.

### Windows

In this example, `ENTRYPOINT` is changed to `cmd.exe`. **CTRL**+**C** is pressed to end the process and stop the container.

console					Copy
<pre>&gt; docker run -it --rm --entrypoint "cmd.exe" myimage</pre>					
Microsoft Windows [Version 10.0.17763.379]					
(c) 2018 Microsoft Corporation. All rights reserved.					
C:\>dir					
Volume in drive C has no label.					
Volume Serial Number is 3005-1E84					
Directory of C:\					
04/09/2019	08:46 AM	<DIR>	app		
03/07/2019	10:25 AM		5,510 License.txt		
04/02/2019	01:35 PM	<DIR>	Program Files		
04/09/2019	01:06 PM	<DIR>	Users		
04/02/2019	01:35 PM	<DIR>	Windows		
			1 File(s)	5,510 bytes	
			4 Dir(s)	21,246,517,248 bytes free	
C:\>^C					

### Linux

In this example, `ENTRYPOINT` is changed to `bash`. The `quit` command is run which ends the process and stop the container.

bash					Copy
------	--	--	--	--	------

```
root@user:~# docker run -it --rm --entrypoint "bash" myimage
root@8515e897c893:/# ls app
myapp.deps.json  myapp.dll  myapp.pdb  myapp.runtimeconfig.json
root@8515e897c893:/# exit
exit
```

## Essential commands


Docker has many different commands that cover what you want to do with your container and images. These Docker commands are essential to managing your containers:

- [docker build](#)
- [docker run](#)
- [docker ps](#)
- [docker stop](#)
- [docker rm](#)
- [docker rmi](#)
- [docker image](#)


## Clean up resources

During this tutorial, you created containers and images. If you want, delete these resources. Use the following commands to


1. List all containers

console	 Copy
<pre>&gt; docker ps -a</pre>	

2. Stop containers that are running. The `CONTAINER_NAME` represents the name automatically assigned to the container.

console	 Copy
<pre>&gt; docker stop CONTAINER_NAME</pre>	

3. Delete the container

console	 Copy
<pre>&gt; docker rm CONTAINER_NAME</pre>	

Next, delete any images that you no longer want on your machine. Delete the image created by your *Dockerfile* and then delete the .NET Core image the *Dockerfile* was based on. You can use the **IMAGE ID** or the **REPOSITORY:TAG** formatted string.

console	 Copy
<pre>docker rmi myimage:latest docker rmi mcr.microsoft.com/dotnet/core/aspnet:3.1</pre>	

Use the `docker images` command to see a list of images installed.

### Note

Image files can be large. Typically, you would remove temporary containers you created while testing and developing your app. You usually keep the base images with the runtime installed if you plan on building other images based on that runtime.

## Next steps

- [Learn how to containerize an ASP.NET Core application.](#)
- [Try the ASP.NET Core Microservice Tutorial.](#)
- [Review the Azure services that support containers.](#)
- [Read about Dockerfile commands.](#)
- [Explore the Container Tools for Visual Studio](#)

---

Is this page helpful?

 Yes  No

---